

▼ CS156 (Introduction to AI), Fall 2022

Homework 10 submission

Roster Name: Preet LNU

Student ID: 014755741

Email address: preet.lnu@sjsu.edu

▼ References and sources

- GAN.MNIST.ipynb

▼ Solution

▼ Load libraries and set random number generator seed

```
import numpy as np
from tensorflow import keras
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Conv2D
from tensorflow.keras.layers import Flatten
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import LeakyReLU
from tensorflow.keras.utils import plot_model
from tensorflow.keras.layers import Reshape
from tensorflow.keras.layers import Conv2DTranspose
from numpy import expand_dims
from numpy import ones
from numpy import zeros
from numpy.random import rand
from numpy.random import randint
from numpy.random import randn
```

```
from numpy import vstack
# ...
np.random.seed(42)
```

▼ Code the solution

```
input_shape = (28, 28, 1)

(x_train, y_train), (x_test, y_test) = keras.datasets.fashion_mnist.load_data()

mnist = np.concatenate([x_train, x_test], axis=0)
mnist = expand_dims(mnist, axis=-1)

mnist = mnist.astype("float32") / 255

mnist.shape

(70000, 28, 28, 1)

def define_discriminator(in_shape=(28,28,1)):
    model = Sequential()
    model.add(Conv2D(64, (3,3), strides=(2, 2), padding='same', input_shape=in_shape))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Dropout(0.4))

    model.add(Conv2D(64, (3,3), strides=(2, 2), padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Dropout(0.4))

    model.add(Conv2D(64, (5,5), strides=(1, 1), padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Dropout(0.4))

    model.add(Flatten())
    model.add(Dense(1, activation='sigmoid'))

    opt = Adam(lr=0.0002, beta_1=0.5)
    model.compile(loss='binary_crossentropy', optimizer=opt, metrics=['accuracy'])
    return model

discriminator = define_discriminator()
discriminator.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 14, 14, 64)	640

leaky_re_lu_7 (LeakyReLU)	(None, 14, 14, 64)	0
dropout_3 (Dropout)	(None, 14, 14, 64)	0
conv2d_5 (Conv2D)	(None, 7, 7, 64)	36928
leaky_re_lu_8 (LeakyReLU)	(None, 7, 7, 64)	0
dropout_4 (Dropout)	(None, 7, 7, 64)	0
conv2d_6 (Conv2D)	(None, 7, 7, 64)	102464
leaky_re_lu_9 (LeakyReLU)	(None, 7, 7, 64)	0
dropout_5 (Dropout)	(None, 7, 7, 64)	0
flatten_1 (Flatten)	(None, 3136)	0
dense_2 (Dense)	(None, 1)	3137

```

=====
Total params: 143,169
Trainable params: 143,169
Non-trainable params: 0

```

```

def define_generator(latent_dim):
    model = Sequential()

    n_nodes = 128 * 7 * 7
    model.add(Dense(n_nodes, input_dim=latent_dim))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Reshape((7, 7, 128)))

    model.add(Conv2DTranspose(128, (4,4), strides=(2,2), padding='same'))
    model.add(LeakyReLU(alpha=0.2))

    model.add(Conv2DTranspose(128, (1,1), strides=(1,1), padding='same'))
    model.add(LeakyReLU(alpha=0.2))

    model.add(Conv2DTranspose(128, (4,4), strides=(2,2), padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Conv2D(1, (7,7), activation='sigmoid', padding='same'))
    return model

latent_dim = 100

generator = define_generator(latent_dim)
generator.summary()

Model: "sequential_4"

```

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 6272)	633472
leaky_re_lu_10 (LeakyReLU)	(None, 6272)	0
reshape_1 (Reshape)	(None, 7, 7, 128)	0
conv2d_transpose_3 (Conv2DTranspose)	(None, 14, 14, 128)	262272
leaky_re_lu_11 (LeakyReLU)	(None, 14, 14, 128)	0
conv2d_transpose_4 (Conv2DTranspose)	(None, 14, 14, 128)	16512
leaky_re_lu_12 (LeakyReLU)	(None, 14, 14, 128)	0
conv2d_transpose_5 (Conv2DTranspose)	(None, 28, 28, 128)	262272
leaky_re_lu_13 (LeakyReLU)	(None, 28, 28, 128)	0
conv2d_7 (Conv2D)	(None, 28, 28, 1)	6273
Total params: 1,180,801		
Trainable params: 1,180,801		
Non-trainable params: 0		

```
def define_gan(g_model, d_model):
```

```
    d_model.trainable = False
```

```
    model = Sequential()
```

```
    model.add(g_model)
```

```
    model.add(d_model)
```

```
    opt = Adam(lr=0.0002, beta_1=0.5)
```

```
    model.compile(loss='binary_crossentropy', optimizer=opt)
```

```
    return model
```

```
gan_model = define_gan(generator, discriminator)
```

```
gan_model.summary()
```

```
Model: "sequential_5"
```

Layer (type)	Output Shape	Param #
sequential_4 (Sequential)	(None, 28, 28, 1)	1180801

```
sequential_3 (Sequential)      (None, 1)      143169
```

```
=====
Total params: 1,323,970
Trainable params: 1,180,801
Non-trainable params: 143,169
=====
```

```
def generate_latent_points(latent_dim, n_samples):
    x_input = randn(latent_dim * n_samples)
    x_input = x_input.reshape(n_samples, latent_dim)
    return x_input

def generate_fake_generator_samples(g_model, latent_dim, n_samples):
    x_input = generate_latent_points(latent_dim, n_samples)
    X = g_model.predict(x_input)
    y = zeros((n_samples, 1))
    return X, y

n_samples = 25
X, _ = generate_fake_generator_samples(generator, latent_dim, n_samples)
for i in range(n_samples):


    plt.subplot(5, 5, 1 + i)

    plt.axis('off')

    plt.imshow(X[i, :, :, 0], cmap='gray_r')

plt.show()
```

```
1/1 [=====] - 0s 274ms/step
```



```
def generate_real_samples(dataset, n_samples):

    ix = randint(0, dataset.shape[0], n_samples)
```

```

X = dataset[ix]
y = ones((n_samples, 1))
return X, y

def generate_fake_samples(g_model, latent_dim, n_samples):
    x_input = generate_latent_points(latent_dim, n_samples)
    X = g_model.predict(x_input)
    y = zeros((n_samples, 1))
    return X, y

def generate_latent_points(latent_dim, n_samples):
    x_input = randn(latent_dim * n_samples)
    x_input = x_input.reshape(n_samples, latent_dim)
    return x_input

def summarize_performance(epoch, g_model, d_model, dataset, latent_dim, n_samples=100)

    X_real, y_real = generate_real_samples(dataset, n_samples)

    _, acc_real = d_model.evaluate(X_real, y_real, verbose=0)

    x_fake, y_fake = generate_fake_samples(g_model, latent_dim, n_samples)

    _, acc_fake = d_model.evaluate(x_fake, y_fake, verbose=0)

    print('>Accuracy real: %.0f%%, fake: %.0f%%' % (acc_real*100, acc_fake*100))

def train(g_model, d_model, gan_model, dataset, latent_dim, n_epochs=10, n_batch=256):
    bat_per_epo = int(dataset.shape[0] / n_batch)
    half_batch = int(n_batch / 2)

    for i in range(n_epochs):

        for j in range(bat_per_epo):

            X_real, y_real = generate_real_samples(dataset, half_batch)

            X_fake, y_fake = generate_fake_samples(g_model, latent_dim, half_batch)

            X, y = vstack((X_real, X_fake)), vstack((y_real, y_fake))

            d_loss, _ = d_model.train_on_batch(X, y)

            X_gan = generate_latent_points(latent_dim, n_batch)

            y_gan = ones((n_batch, 1))

            g_loss = gan_model.train_on_batch(X_gan, y_gan)

            print('>%d, %d/%d, d_loss=%.3f, g_loss=%.3f' % (i+1, j+1, bat_per_epo, d_

```

```
summarize_performance(i, g_model, d_model, dataset, latent_dim)
```

```
return g_model
```

```
latent_dim = 100
```

```
trained_generator = train(generator, discriminator, gan_model, mnist, latent_dim, 10)
```

```

*/* [-----] - 1s 202ms/step
[>] >10, 246/273, d_loss=0.683, g_loss=0.693
4/4 [=====] - 1s 216ms/step
>10, 247/273, d_loss=0.692, g_loss=0.719
4/4 [=====] - 1s 227ms/step
>10, 248/273, d_loss=0.671, g_loss=0.773
4/4 [=====] - 1s 217ms/step
>10, 249/273, d_loss=0.675, g_loss=0.777
4/4 [=====] - 1s 213ms/step
>10, 250/273, d_loss=0.688, g_loss=0.778
4/4 [=====] - 1s 240ms/step
>10, 251/273, d_loss=0.669, g_loss=0.722
4/4 [=====] - 1s 219ms/step
>10, 252/273, d_loss=0.670, g_loss=0.704
4/4 [=====] - 1s 208ms/step
>10, 253/273, d_loss=0.668, g_loss=0.731
4/4 [=====] - 1s 214ms/step
>10, 254/273, d_loss=0.674, g_loss=0.754
4/4 [=====] - 1s 213ms/step
>10, 255/273, d_loss=0.668, g_loss=0.723
4/4 [=====] - 1s 215ms/step
>10, 256/273, d_loss=0.672, g_loss=0.697
4/4 [=====] - 1s 210ms/step
>10, 257/273, d_loss=0.675, g_loss=0.673
4/4 [=====] - 1s 221ms/step
>10, 258/273, d_loss=0.663, g_loss=0.675
4/4 [=====] - 1s 221ms/step
>10, 259/273, d_loss=0.701, g_loss=0.764
4/4 [=====] - 1s 216ms/step
>10, 260/273, d_loss=0.688, g_loss=0.876
4/4 [=====] - 1s 218ms/step

>10, 261/273, d_loss=0.680, g_loss=0.909
4/4 [=====] - 1s 214ms/step
>10, 262/273, d_loss=0.689, g_loss=0.777
4/4 [=====] - 1s 217ms/step
>10, 263/273, d_loss=0.681, g_loss=0.688
4/4 [=====] - 1s 211ms/step
>10, 264/273, d_loss=0.669, g_loss=0.624
4/4 [=====] - 1s 218ms/step
>10, 265/273, d_loss=0.678, g_loss=0.635
4/4 [=====] - 1s 218ms/step
>10, 266/273, d_loss=0.654, g_loss=0.669
4/4 [=====] - 1s 215ms/step
>10, 267/273, d_loss=0.675, g_loss=0.737
4/4 [=====] - 1s 220ms/step
>10, 268/273, d_loss=0.684, g_loss=0.843
4/4 [=====] - 1s 220ms/step

```

```
>10, 269/273, d_loss=0.671, g_loss=0.885
4/4 [=====] - 1s 227ms/step
>10, 270/273, d_loss=0.678, g_loss=0.872
4/4 [=====] - 1s 220ms/step
>10, 271/273, d_loss=0.651, g_loss=0.767
4/4 [=====] - 1s 217ms/step
>10, 272/273, d_loss=0.658, g_loss=0.769
4/4 [=====] - 1s 219ms/step
>10, 273/273, d_loss=0.676, g_loss=0.720
4/4 [=====] - 1s 161ms/step
>Accuracy real: 76%, fake: 39%
```

```
def generate_latent_points(latent_dim, n_samples):
```

```
    x_input = randn(latent_dim * n_samples)
```

```
    x_input = x_input.reshape(n_samples, latent_dim)
```

```
    return x_input
```

```
def display_plot(examples, n):
```

```
    for i in range(n * n):
```

```
        plt.subplot(n, n, 1 + i)
```

```
        plt.axis('off')
```

```
        plt.imshow(examples[i, :, :, 0], cmap='gray_r')
```

```
    plt.show()
```

```
latent_points = generate_latent_points(100, 25)
```

```
X = trained_generator.predict(latent_points)
```

```
display_plot(X, 5)
```

```
1/1 [=====] - 0s 180ms/step
```



[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 5:30 AM

